

FinalSVM

January 10, 2020

```
[18]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import os

from sklearn.svm import SVC
from sklearn.datasets import fetch_openml

mnist = fetch_openml('mnist_784', version=1, cache=True)

X = pd.DataFrame(mnist.data)
Y= pd.Series(mnist.target).astype('int').astype('category')
```

```
[19]: cleanTestX = X.tail(2000)
cleanTestY = Y.tail(2000)

X = X.head(10000)
Y = Y.head(10000)
```

```
[20]: print("X data")
print(X.shape)
# keep in mind 28 by 28 pixel values
print ("Y data")
print(Y.shape)
#y / output only needs one colum saying which class the image fell into
```

```
X data
(10000, 784)
Y data
(10000,)
```

```
[21]: from sklearn.model_selection import train_test_split
X_train, X_val, y_train, y_val = train_test_split(X, Y, test_size = 0.20)
```

```
[22]: ## Testing out linear kernel
svcLinear = SVC(kernel= 'linear')
svcLinear.fit(X_train,y_train)
```

```

predicted_Linear =svcLinear.predict(cleanTestX)
from sklearn.metrics import classification_report, accuracy_score

print(classification_report(cleanTestY,predicted_Linear))
accuracy_score(cleanTestY, predicted_Linear)

```

	precision	recall	f1-score	support
0	0.95	0.97	0.96	207
1	0.98	0.99	0.98	230
2	0.92	0.88	0.90	198
3	0.93	0.90	0.91	207
4	0.91	0.95	0.93	194
5	0.90	0.91	0.91	169
6	0.93	0.96	0.94	202
7	0.94	0.94	0.94	215
8	0.90	0.91	0.91	187
9	0.93	0.87	0.90	191
accuracy			0.93	2000
macro avg	0.93	0.93	0.93	2000
weighted avg	0.93	0.93	0.93	2000

[22]: 0.9305

```

[23]: ## Testing Polynomial Kernel
svcPoly = SVC(kernel= 'poly' )
svcPoly.fit(X_train,y_train)

predicted_Poly =svcPoly.predict(cleanTestX)

print(classification_report(cleanTestY,predicted_Poly))
accuracy_score(cleanTestY, predicted_Poly)

```

	precision	recall	f1-score	support
0	0.97	0.98	0.97	207
1	0.99	1.00	0.99	230
2	0.96	0.93	0.94	198
3	0.99	0.97	0.98	207
4	0.95	0.98	0.97	194
5	0.94	0.96	0.95	169
6	0.96	0.99	0.97	202
7	0.97	0.97	0.97	215
8	0.96	0.95	0.95	187
9	0.99	0.95	0.97	191

accuracy			0.97	2000
macro avg	0.97	0.97	0.97	2000
weighted avg	0.97	0.97	0.97	2000

[23]: 0.968

```
[24]: ## Testing RBF kernel
svcRBF = SVC(kernel= 'rbf' )
svcRBF.fit(X_train,y_train)

predicted_RBF =svcRBF.predict(cleanTestX)

print(classification_report(cleanTestY,predicted_RBF))
accuracy_score(cleanTestY, predicted_RBF)
```

	precision	recall	f1-score	support
0	0.97	0.99	0.98	207
1	1.00	1.00	1.00	230
2	0.95	0.94	0.95	198
3	0.99	0.98	0.99	207
4	0.96	0.97	0.97	194
5	0.99	0.95	0.97	169
6	0.95	0.99	0.97	202
7	0.98	0.97	0.98	215
8	0.95	0.97	0.96	187
9	0.98	0.96	0.97	191
accuracy			0.97	2000
macro avg	0.97	0.97	0.97	2000
weighted avg	0.97	0.97	0.97	2000

[24]: 0.9725

```
[25]: ## Testing Sigmoid kernel
svcSigmoid = SVC(kernel= 'sigmoid' )
svcSigmoid.fit(X_train,y_train)

predicted_Sigmoid =svcSigmoid.predict(cleanTestX)

print(classification_report(cleanTestY,predicted_Sigmoid))
accuracy_score(cleanTestY, predicted_Sigmoid)
```

precision	recall	f1-score	support
-----------	--------	----------	---------

0	0.83	0.89	0.86	207
1	0.95	0.97	0.96	230
2	0.81	0.79	0.80	198
3	0.88	0.82	0.84	207
4	0.88	0.94	0.91	194
5	0.72	0.77	0.74	169
6	0.94	0.88	0.91	202
7	0.93	0.92	0.93	215
8	0.78	0.73	0.75	187
9	0.87	0.85	0.86	191
accuracy			0.86	2000
macro avg		0.86	0.86	2000
weighted avg		0.86	0.86	2000

[25]: 0.861

[26]: *### NOW TO FIND A MORE STRUCTURED APPROACH TO TUNING BY MAKING USE OF A GRID_{SEARCH}*

[27]: `from sklearn.model_selection import GridSearchCV`

[]:

```
# defining parameter range
param_grid = {'C': [1,2,3],
              'gamma': [0.0000002, 0.0000003, 0.0000004, 0.00000044],
              'kernel': ['rbf']}

grid = GridSearchCV(SVC(), param_grid, refit = True, verbose = 3)

grid.fit(X_train, y_train)

predicted_yNew =grid.predict(X_val)
print(accuracy_score(y_val, predicted_yNew))
print("***** ----- *****")
print(classification_report(y_val,predicted_yNew))

print(grid.best_estimator_)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

[CV] C=1, gamma=2e-07, kernel=rbf ...

[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] ... C=1, gamma=2e-07, kernel=rbf, score=0.952, total= 18.0s

[CV] C=1, gamma=2e-07, kernel=rbf ...

```
[Parallel(n_jobs=1)]: Done    1 out of    1 | elapsed:  18.0s remaining:  0.0s
[CV] ... C=1, gamma=2e-07, kernel=rbf, score=0.961, total= 17.9s
[CV] C=1, gamma=2e-07, kernel=rbf ...
[Parallel(n_jobs=1)]: Done    2 out of    2 | elapsed:  35.9s remaining:  0.0s
[CV] ... C=1, gamma=2e-07, kernel=rbf, score=0.963, total= 18.0s
[CV] C=1, gamma=2e-07, kernel=rbf ...
[CV] ... C=1, gamma=2e-07, kernel=rbf, score=0.953, total= 18.2s
[CV] C=1, gamma=2e-07, kernel=rbf ...
[CV] ... C=1, gamma=2e-07, kernel=rbf, score=0.954, total= 18.4s
[CV] C=1, gamma=3e-07, kernel=rbf ...
[CV] ... C=1, gamma=3e-07, kernel=rbf, score=0.956, total= 22.1s
[CV] C=1, gamma=3e-07, kernel=rbf ...
[CV] ... C=1, gamma=3e-07, kernel=rbf, score=0.966, total= 22.5s
[CV] C=1, gamma=3e-07, kernel=rbf ...
[CV] ... C=1, gamma=3e-07, kernel=rbf, score=0.969, total= 27.9s
[CV] C=1, gamma=3e-07, kernel=rbf ...
[CV] ... C=1, gamma=3e-07, kernel=rbf, score=0.960, total= 23.0s
[CV] C=1, gamma=3e-07, kernel=rbf ...
[CV] ... C=1, gamma=3e-07, kernel=rbf, score=0.960, total= 21.8s
[CV] C=1, gamma=4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4e-07, kernel=rbf, score=0.959, total= 26.5s
[CV] C=1, gamma=4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4e-07, kernel=rbf, score=0.969, total= 26.5s
[CV] C=1, gamma=4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4e-07, kernel=rbf, score=0.973, total= 26.4s
[CV] C=1, gamma=4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4e-07, kernel=rbf, score=0.960, total= 27.5s
[CV] C=1, gamma=4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4e-07, kernel=rbf, score=0.961, total= 27.9s
[CV] C=1, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4.4e-07, kernel=rbf, score=0.960, total= 29.8s
[CV] C=1, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4.4e-07, kernel=rbf, score=0.971, total= 31.0s
[CV] C=1, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4.4e-07, kernel=rbf, score=0.973, total= 31.8s
[CV] C=1, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4.4e-07, kernel=rbf, score=0.961, total= 31.1s
[CV] C=1, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=1, gamma=4.4e-07, kernel=rbf, score=0.962, total= 30.1s
[CV] C=2, gamma=2e-07, kernel=rbf ...
[CV] ... C=2, gamma=2e-07, kernel=rbf, score=0.958, total= 18.3s
[CV] C=2, gamma=2e-07, kernel=rbf ...
[CV] ... C=2, gamma=2e-07, kernel=rbf, score=0.968, total= 18.2s
[CV] C=2, gamma=2e-07, kernel=rbf ...
[CV] ... C=2, gamma=2e-07, kernel=rbf, score=0.966, total= 18.8s
[CV] C=2, gamma=2e-07, kernel=rbf ...
```

```
[CV] ... C=2, gamma=2e-07, kernel=rbf, score=0.958, total= 18.6s
[CV] C=2, gamma=2e-07, kernel=rbf ...
[CV] ... C=2, gamma=2e-07, kernel=rbf, score=0.959, total= 18.6s
[CV] C=2, gamma=3e-07, kernel=rbf ...
[CV] ... C=2, gamma=3e-07, kernel=rbf, score=0.961, total= 22.6s
[CV] C=2, gamma=3e-07, kernel=rbf ...
[CV] ... C=2, gamma=3e-07, kernel=rbf, score=0.971, total= 22.7s
[CV] C=2, gamma=3e-07, kernel=rbf ...
[CV] ... C=2, gamma=3e-07, kernel=rbf, score=0.969, total= 22.5s
[CV] C=2, gamma=3e-07, kernel=rbf ...
[CV] ... C=2, gamma=3e-07, kernel=rbf, score=0.960, total= 22.9s
[CV] C=2, gamma=3e-07, kernel=rbf ...
[CV] ... C=2, gamma=3e-07, kernel=rbf, score=0.961, total= 22.8s
[CV] C=2, gamma=4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4e-07, kernel=rbf, score=0.962, total= 27.0s
[CV] C=2, gamma=4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4e-07, kernel=rbf, score=0.973, total= 27.2s
[CV] C=2, gamma=4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4e-07, kernel=rbf, score=0.974, total= 28.8s
[CV] C=2, gamma=4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4e-07, kernel=rbf, score=0.963, total= 28.2s
[CV] C=2, gamma=4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4e-07, kernel=rbf, score=0.961, total= 28.2s
[CV] C=2, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4.4e-07, kernel=rbf, score=0.964, total= 30.9s
[CV] C=2, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4.4e-07, kernel=rbf, score=0.975, total= 29.8s
[CV] C=2, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4.4e-07, kernel=rbf, score=0.974, total= 29.0s
[CV] C=2, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4.4e-07, kernel=rbf, score=0.962, total= 28.5s
[CV] C=2, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=2, gamma=4.4e-07, kernel=rbf, score=0.961, total= 28.9s
[CV] C=3, gamma=2e-07, kernel=rbf ...
[CV] ... C=3, gamma=2e-07, kernel=rbf, score=0.961, total= 17.3s
[CV] C=3, gamma=2e-07, kernel=rbf ...
[CV] ... C=3, gamma=2e-07, kernel=rbf, score=0.969, total= 18.3s
[CV] C=3, gamma=2e-07, kernel=rbf ...
[CV] ... C=3, gamma=2e-07, kernel=rbf, score=0.965, total= 17.4s
[CV] C=3, gamma=2e-07, kernel=rbf ...
[CV] ... C=3, gamma=2e-07, kernel=rbf, score=0.960, total= 17.1s
[CV] C=3, gamma=2e-07, kernel=rbf ...
[CV] ... C=3, gamma=2e-07, kernel=rbf, score=0.960, total= 17.4s
[CV] C=3, gamma=3e-07, kernel=rbf ...
[CV] ... C=3, gamma=3e-07, kernel=rbf, score=0.961, total= 20.9s
[CV] C=3, gamma=3e-07, kernel=rbf ...
[CV] ... C=3, gamma=3e-07, kernel=rbf, score=0.971, total= 20.9s
[CV] C=3, gamma=3e-07, kernel=rbf ...
```

```

[CV] ... C=3, gamma=3e-07, kernel=rbf, score=0.971, total= 21.1s
[CV] C=3, gamma=3e-07, kernel=rbf ...
[CV] ... C=3, gamma=3e-07, kernel=rbf, score=0.961, total= 21.1s
[CV] C=3, gamma=3e-07, kernel=rbf ...
[CV] ... C=3, gamma=3e-07, kernel=rbf, score=0.962, total= 22.0s
[CV] C=3, gamma=4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4e-07, kernel=rbf, score=0.964, total= 26.6s
[CV] C=3, gamma=4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4e-07, kernel=rbf, score=0.973, total= 27.4s
[CV] C=3, gamma=4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4e-07, kernel=rbf, score=0.974, total= 26.4s
[CV] C=3, gamma=4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4e-07, kernel=rbf, score=0.963, total= 26.2s
[CV] C=3, gamma=4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4e-07, kernel=rbf, score=0.961, total= 26.0s
[CV] C=3, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4.4e-07, kernel=rbf, score=0.964, total= 28.1s
[CV] C=3, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4.4e-07, kernel=rbf, score=0.975, total= 29.1s
[CV] C=3, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4.4e-07, kernel=rbf, score=0.973, total= 29.1s
[CV] C=3, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4.4e-07, kernel=rbf, score=0.962, total= 29.3s
[CV] C=3, gamma=4.4e-07, kernel=rbf ...
[CV] ... C=3, gamma=4.4e-07, kernel=rbf, score=0.961, total= 28.4s

```

[Parallel(n_jobs=1)]: Done 60 out of 60 | elapsed: 24.3min finished

0.975

***** ----- *****

	precision	recall	f1-score	support
0	0.98	0.99	0.99	187
1	1.00	0.98	0.99	219
2	0.98	0.99	0.99	202
3	0.98	0.96	0.97	228
4	0.98	0.97	0.98	195
5	0.96	0.99	0.97	171
6	0.97	0.96	0.97	199
7	0.95	0.98	0.97	207
8	0.98	0.98	0.98	201
9	0.95	0.95	0.95	191
accuracy			0.97	2000
macro avg	0.97	0.98	0.97	2000
weighted avg	0.98	0.97	0.98	2000

SVC(C=2, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
 decision_function_shape='ovr', degree=3, gamma=4.4e-07, kernel='rbf',

```
max_iter=-1, probability=False, random_state=None, shrinking=True,  
tol=0.001, verbose=False)
```

```
[29]: # defining parameter range  
param_gridPoly = {'degree': [2,3,4,5],  
                  'kernel': ['poly']}  
  
gridP = GridSearchCV(SVC(), param_gridPoly, refit = True, verbose = 3)  
  
gridP.fit(X_train, y_train)  
  
predicted_yP = gridP.predict(X_val)  
print(accuracy_score(y_val, predicted_yP))  
print("***** ----- *****")  
print(classification_report(y_val,predicted_yP))  
  
print(gridP.best_estimator_)
```

```
Fitting 5 folds for each of 4 candidates, totalling 20 fits  
[CV] degree=2, kernel=poly ...  
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[CV] ... degree=2, kernel=poly, score=0.951, total= 14.6s  
[CV] degree=2, kernel=poly ...  
[Parallel(n_jobs=1)]: Done 1 out of 1 | elapsed: 14.6s remaining: 0.0s  
[CV] ... degree=2, kernel=poly, score=0.956, total= 14.6s  
[CV] degree=2, kernel=poly ...  
[Parallel(n_jobs=1)]: Done 2 out of 2 | elapsed: 29.2s remaining: 0.0s  
[CV] ... degree=2, kernel=poly, score=0.956, total= 14.6s  
[CV] degree=2, kernel=poly ...  
[CV] ... degree=2, kernel=poly, score=0.949, total= 14.6s  
[CV] degree=2, kernel=poly ...  
[CV] ... degree=2, kernel=poly, score=0.953, total= 14.7s  
[CV] degree=3, kernel=poly ...  
[CV] ... degree=3, kernel=poly, score=0.939, total= 17.2s  
[CV] degree=3, kernel=poly ...  
[CV] ... degree=3, kernel=poly, score=0.953, total= 17.1s  
[CV] degree=3, kernel=poly ...  
[CV] ... degree=3, kernel=poly, score=0.950, total= 17.1s  
[CV] degree=3, kernel=poly ...  
[CV] ... degree=3, kernel=poly, score=0.946, total= 17.1s  
[CV] degree=3, kernel=poly ...  
[CV] ... degree=3, kernel=poly, score=0.950, total= 17.2s  
[CV] degree=4, kernel=poly ...  
[CV] ... degree=4, kernel=poly, score=0.924, total= 20.6s  
[CV] degree=4, kernel=poly ...
```

```

[CV] ... degree=4, kernel=poly, score=0.917, total= 20.6s
[CV] degree=4, kernel=poly ...
[CV] ... degree=4, kernel=poly, score=0.920, total= 20.5s
[CV] degree=4, kernel=poly ...
[CV] ... degree=4, kernel=poly, score=0.920, total= 20.8s
[CV] degree=4, kernel=poly ...
[CV] ... degree=4, kernel=poly, score=0.935, total= 20.6s
[CV] degree=5, kernel=poly ...
[CV] ... degree=5, kernel=poly, score=0.866, total= 24.4s
[CV] degree=5, kernel=poly ...
[CV] ... degree=5, kernel=poly, score=0.875, total= 24.4s
[CV] degree=5, kernel=poly ...
[CV] ... degree=5, kernel=poly, score=0.874, total= 25.5s
[CV] degree=5, kernel=poly ...
[CV] ... degree=5, kernel=poly, score=0.887, total= 25.5s
[CV] degree=5, kernel=poly ...
[CV] ... degree=5, kernel=poly, score=0.884, total= 24.2s

[Parallel(n_jobs=1)]: Done 20 out of 20 | elapsed: 6.4min finished
0.955
***** ----- *****

```

	precision	recall	f1-score	support
0	0.96	0.98	0.97	187
1	0.96	0.98	0.97	219
2	0.98	0.97	0.97	202
3	0.96	0.94	0.95	228
4	0.94	0.94	0.94	195
5	0.95	0.95	0.95	171
6	0.96	0.95	0.96	199
7	0.96	0.96	0.96	207
8	0.96	0.94	0.95	201
9	0.90	0.93	0.92	191
accuracy			0.95	2000
macro avg	0.95	0.95	0.95	2000
weighted avg	0.96	0.95	0.96	2000

```

SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=2, gamma='scale', kernel='poly',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)

```

```

[30]: #On Test data not on validation
predicted_yPFinal =gridP.predict(cleanTestX)
print("Poly Trained results \n"
      +"classification_report(cleanTestY,predicted_yPFinal))
```

Poly Trained results

	precision	recall	f1-score	support
0	0.97	0.99	0.98	207
1	1.00	1.00	1.00	230
2	0.95	0.94	0.94	198
3	1.00	0.97	0.98	207
4	0.97	0.98	0.98	194
5	0.96	0.95	0.96	169
6	0.96	0.99	0.97	202
7	0.98	0.97	0.97	215
8	0.95	0.95	0.95	187
9	0.98	0.95	0.97	191
accuracy			0.97	2000
macro avg	0.97	0.97	0.97	2000
weighted avg	0.97	0.97	0.97	2000

```
[31]: #On Test data not on validation
predicted_yRBGFinal =grid.predict(cleanTestX)
print("RBG Trained results \n")
print(classification_report(cleanTestY,predicted_yRBGFinal))
```

RBG Trained results

	precision	recall	f1-score	support
0	0.98	0.99	0.98	207
1	1.00	1.00	1.00	230
2	0.95	0.97	0.96	198
3	1.00	0.98	0.99	207
4	0.96	0.97	0.97	194
5	0.99	0.96	0.97	169
6	0.97	0.99	0.98	202
7	0.98	0.97	0.97	215
8	0.96	0.97	0.97	187
9	0.98	0.96	0.97	191
accuracy			0.98	2000
macro avg	0.98	0.98	0.98	2000
weighted avg	0.98	0.98	0.98	2000

[]: